

Multivariate Bayesian P-spline estimation of spectral density matrices, with application to LISA TDI noise

Avi Vajpeyi¹, Renate Meyer¹, Patricio Maturana-Russel^{1,2}, and Jianan Liu¹

¹*Department of Statistics, The University of Auckland, Auckland, New Zealand and*

²*Department of Mathematical Sciences, Auckland University of Technology, Auckland, New Zealand*

(Dated: May 15, 2026)

We present a Bayesian multivariate P-spline method for estimating the frequency-dependent cross-spectral density matrix of stationary vector time series. We model the lower-triangular and diagonal components of the frequency-varying Cholesky decomposition of the inverse spectral density matrix, enforcing Hermitian positive definiteness at every frequency. Model flexibility and smoothness for each Cholesky component is induced by penalised B-spline priors. Inference proceeds under a blocked, coarse-grained Whittle likelihood with safe-Bayes η -tempering to stabilise posterior calibration, and is performed with the No-U-Turn Sampler initialised from draws of a stochastic variational approximation. Each likelihood evaluation scales linearly with the number of frequencies and basis functions, and quadratically with the number of channels. On synthetic VAR(2) benchmarks with known ground truth, the method accurately recovers diagonal and cross-spectral structure, attains near-nominal coverage, and achieves low relative integrated absolute error (RIAE) that improves with sample size and is robust across a broad range of spline basis sizes. We illustrate the relevance to LISA by estimating the multichannel noise spectral matrix for time-delay interferometry combinations, and demonstrate that AET-restricted models fail on realistic asymmetric LISA noise by an order of magnitude in relative integrated absolute error.

I. INTRODUCTION

Accurate estimation of the noise power spectral density (PSD) underpins gravitational-wave (GW) data analysis: search, parameter estimation, and stochastic-background inference all require a reliable noise model across the detector band. For LIGO and Virgo, parametric models augmented with spline components, most notably BayesLine [1], are widely used. The space-based detector LISA [2] poses a harder problem: its three time-delay interferometry (TDI) channels are correlated, and their noise properties vary smoothly but non-trivially across the millihertz band. A single-channel noise model cannot capture the cross-spectral correlations that matter for joint inference.

Nonparametric/flexible Bayesian approaches to PSD estimation have a well-developed statistical foundation, characterised by an expansion of the PSD in terms of an infinite/large number of basis functions combined with noninformative priors on the coefficients. Examples include Bernstein polynomials with weights induced by Dirichlet process priors [3], Akima splines with Beta distributions on control frequencies [4], cosines with a discounted regularized horseshoe prior [5], and B-splines with knots and weights induced by Dirichlet process priors [6]. With an a priori unknown and potentially infinite number of basis functions, posterior computation requires trans-dimensional Markov chain Monte Carlo (MCMC) algorithms, which are computationally expensive. Penalised B-splines (P-splines) [7] avoid the trans-dimensional step by representing smooth functions as linear combinations of a finite but large number of B-spline basis functions, with a roughness penalty on adjacent coefficients that prevents overfitting. In a Bayesian

formulation this penalty can be incorporated naturally into the prior distribution of the spline coefficients. Applied to spectral density estimation via the Whittle likelihood, this yields a smooth, data-driven estimator that requires no parametric noise template. Univariate Bayesian P-spline PSD methods were introduced by Maturana-Russel and Meyer [8] and have since been applied to GW data, including recent LISA-specific work that introduced log-spline representations tuned to the steep low-frequency rise of LISA noise [9], as well as nonparametric stochastic-background searches that jointly model unknown noise shapes [10, 11].

These univariate approaches model each channel independently. This is justified in the idealised equal-arm stationary LISA configuration, which has vanishing off-diagonal cross-spectra between the three TDI channels, but realistic noise (unequal arm lengths, per-MOSA imbalance, and instrumental breathing [12]) generates non-trivial off-diagonal entries in the 3×3 spectral density matrix $\mathbf{S}(f)$, whose normalised magnitudes (the channel-pair coherences) quantify the cross-channel correlation that single-channel models discard. Treating these cross-spectra as zero when they are not can degrade sensitivity to stochastic GW backgrounds and bias parameter estimation [13, 14]. Any extension of univariate methods to nonparametrically model $\mathbf{S}(f)$ needs to ensure Hermitian positive definiteness at every frequency while allowing each matrix element to vary smoothly with frequency.

Existing nonparametric treatments of LISA TDI spectral matrices take two main routes. Muratore, Gair and Speri [15] model the diagonal PSDs and the real and imaginary parts of the cross-spectra of $\mathbf{S}(f)$ separately using natural cubic splines. This is flexible, but does not by itself guarantee that the assembled matrix is pos-

itive definite. A more structured alternative starts from the six single-link measurements, assumes equal and uncorrelated test-mass acceleration and optical-metrology-system noises with univariate spectra S_{TM} and S_{OMS} , and approximates the delay operators defining the XYZ TDI variables as complex phase factors. With fixed, known TDI transfer functions, the resulting XYZ spectral matrix is Hermitian positive definite while only S_{TM} and S_{OMS} are modelled nonparametrically with cubic splines [4].

In this paper, we model the PSD matrix of the XYZ TDI variables directly by extending the Bayesian P-spline PSD framework [8] to the multivariate p -dimensional setting. We parametrise $\mathbf{S}(f)^{-1}$ through its Cholesky decomposition [16, 17], which factorises the posterior density into p independent components that can be sampled in parallel and guarantees the Hermitian positive definiteness condition. Each block's log-diagonal spectra and real and imaginary parts of the complex off-diagonal cross-spectra are modeled with penalised B-splines under a hierarchical smoothing prior. This prior is updated by a blocked, coarse-grained Whittle likelihood with safe-Bayes η -tempering [18] to stabilise posterior calibration. Each chain is initialised from a draw of a stochastic variational approximation [19], with sampling performed by the No-U-Turn Sampler [20]. We validate the method on a bivariate VAR(2) benchmark against Liu *et al.* [21], a three-channel VAR(2) simulation, and a simulated LISA data segment.

The remainder of the paper is organised as follows. Section II develops the statistical model and the inference scheme. Section III presents simulation studies on a 3D VAR(2) benchmark and on simulated LISA TDI data. Section IV summarises the results, discusses limitations and future directions.

II. MULTIVARIATE P-SPLINES

A. Likelihood

A glossary of the symbols used throughout this paper is provided in Table III (Appendix A). Readers may find it useful to keep this table open while reading the rest of this section. Let $\mathbf{Z} = (\mathbf{Z}_1, \dots, \mathbf{Z}_n)^\top \in \mathbb{R}^{n \times p}$ be a p -dimensional stationary, mean-zero time series sampled at intervals $\Delta_t = 1/(2f_{Ny})$, so that $\mathbf{Z}_\nu = \mathbf{Z}(\nu\Delta_t)$ for $\nu = 1, \dots, n$, where f_{Ny} is the Nyquist frequency. The total observation time is $T = n\Delta_t$ and the frequency resolution is

$$\Delta_f = \frac{1}{n\Delta_t} = \frac{1}{T}. \quad (1)$$

The discrete Fourier transform (DFT) of $\mathbf{Z}$ is

$$\mathbf{d}(f_k) = \Delta_t \sum_{t=1}^n \mathbf{Z}_t \exp\left(-2\pi i \frac{k}{n} t\right), \quad (2)$$

with $f_k = k/T$ for $k = 1, \dots, N$, where $N = \lfloor n/2 \rfloor$.

For stationary series with absolutely summable matrix autocovariances $\sum_{h=-\infty}^{\infty} \|\Gamma(h)\| < \infty$, the Fourier coefficients $\mathbf{d}(f_k)$ are asymptotically independent, complex Gaussian with mean zero and covariance $T\mathbf{S}(f_k)$, where

$$\mathbf{S}(f) = \frac{1}{2f_{Ny}} \sum_{\nu=-\infty}^{\infty} \Gamma(\nu\Delta_t) \exp(-2\pi i f \nu\Delta_t)$$

is the two-sided spectral density matrix — the Fourier transform of the autocovariance $\Gamma(h) = \mathbb{E}(\mathbf{Z}_t \mathbf{Z}_{t+h}^\top)$. This asymptotic Gaussian approximation is the basis of the multivariate Whittle likelihood [22]

$$\mathcal{L}(\mathbf{d}|\mathbf{S}) \propto \prod_{k=1}^N \det(\mathbf{S}(f_k))^{-1} \times \exp\left(-\frac{1}{T} \mathbf{d}(f_k)^* \mathbf{S}(f_k)^{-1} \mathbf{d}(f_k)\right), \quad (3)$$

where $\mathbf{d}(f_k)^*$ denotes the conjugate transpose of $\mathbf{d}(f_k)$ and $\mathbf{S}(f_k)$ is a Hermitian positive semidefinite matrix.

The diagonal entries $S_{ii}(f)$ are the power spectral densities (PSDs) of channel i , while the off-diagonals $S_{ij}(f)$ for $i \neq j$ are the (complex) cross-spectra. A scale-free summary of the linear dependence between channels i and j at frequency f is the complex coherence,

$$C_{ij}(f) = \frac{S_{ij}(f)}{\sqrt{S_{ii}(f) S_{jj}(f)}}, \quad |C_{ij}(f)| \in [0, 1], \quad (4)$$

whose magnitude is zero when channels i and j are uncorrelated at frequency f and unity when they are perfectly linearly related. We use $|C_{ij}(f)|$ to visualise off-diagonal recovery in Sec. III.

Blocking. To trade frequency resolution for periodogram stability we partition $\mathbf{Z}$ into N_b equal, non-overlapping blocks $(\mathbf{Z}^{(1)}, \dots, \mathbf{Z}^{(N_b)})$, each of duration $T_b = T/N_b$ and containing n/N_b samples (illustrated in Fig. 1)¹. Denoting the DFT of block i by $\mathbf{d}^{(i)}$, and assuming blocks are approximately independent under stationarity, the blocked Whittle likelihood factorises as

$$\mathcal{L}_b(\mathbf{d}|\mathbf{S}) = \prod_{i=1}^{N_b} \mathcal{L}(\mathbf{d}^{(i)}|\mathbf{S}). \quad (5)$$

Collecting the sum over blocks into a single Hermitian matrix

$$\mathbf{Y}(f_k) = \sum_{i=1}^{N_b} \mathbf{I}^{(i)}(f_k) = N_b \bar{\mathbf{I}}(f_k), \quad (6)$$

¹ Each block may be tapered by a window function w_t before the DFT to suppress spectral leakage. In that case the block periodogram is rescaled by the inverse effective noise bandwidth $\text{ENBW} = n_b \sum_t w_t^2 / (\sum_t w_t)^2$ (with $n_b = n/N_b$), so that the complex-Wishart mean of $\mathbf{Y}(f_k)$ remains $T_b \mathbf{S}(f_k)$ and Eqs. (7)–(10) apply unchanged.

where $\mathbf{I}^{(i)}(f_k) = \mathbf{d}^{(i)}(f_k)\mathbf{d}^{(i)}(f_k)^*$ is the block- i periodogram and $\bar{\mathbf{I}}(f) = N_b^{-1} \sum_{i=1}^{N_b} \mathbf{I}^{(i)}(f)$ is the block-averaged periodogram (the Welch estimator), and applying the trace cyclic property $\text{tr}(ABC) = \text{tr}(CAB)$, the blocked likelihood reduces to

$$\mathcal{L}_b(\mathbf{d}|\mathbf{S}) \propto \prod_{k=1}^{n/(2N_b)} |\mathbf{S}(f_k)|^{-N_b} \exp\left(-\frac{1}{T_b} \text{tr}[\mathbf{S}(f_k)^{-1} \mathbf{Y}(f_k)]\right). \quad (7)$$

Under the asymptotic Gaussian approximation, $\mathbf{Y}(f_k)$ follows a complex Wishart distribution² from summing N_b assumed-independent single-block periodogram contributions:

$$\mathbf{I}^{(i)}(f_k) \sim \mathcal{CW}_d(T_b \mathbf{S}(f_k), 1), \quad (8)$$

$$\mathbf{Y}(f_k) \sim \mathcal{CW}_d(T_b \mathbf{S}(f_k), N_b). \quad (9)$$

Eigendecomposition. Because $\mathbf{Y}(f_k)$ is Hermitian positive semidefinite, it admits an eigendecomposition $\mathbf{Y}(f_k) = \sum_{\nu=1}^p \lambda_{\nu}^{(k)} \mathbf{v}_{\nu}^{(k)} \mathbf{v}_{\nu}^{(k)*} = \sum_{\nu=1}^p \mathbf{u}_{\nu}^{(k)} \mathbf{u}_{\nu}^{(k)*}$, with rescaled eigenvectors $\mathbf{u}_{\nu}^{(k)} = \sqrt{\lambda_{\nu}^{(k)}} \mathbf{v}_{\nu}^{(k)}$. Substituting into the trace in (7) gives the eigenvector form of the blocked likelihood,

$$\mathcal{L}_b(\mathbf{d}|\mathbf{S}) \propto \prod_{k=1}^{n/(2N_b)} |\mathbf{S}(f_k)|^{-N_b} \times \exp\left(-\frac{1}{T_b} \sum_{\nu=1}^p \mathbf{u}_{\nu}^{(k)*} \mathbf{S}(f_k)^{-1} \mathbf{u}_{\nu}^{(k)}\right). \quad (10)$$

This follows from the spectral decomposition of the Hermitian positive semidefinite $\mathbf{Y}(f_k)$ and the linearity of the trace. The eigenvector representation is used computationally because it turns the trace contribution into a sum of quadratic forms in the Cholesky parameters, which is the form that factorises into the per-channel regressions below.

a. Coarse graining. Blocking reduces the number of periodogram matrices from N to $n/(2N_b)$, but each of the remaining frequencies still requires a separate evaluation of $\mathbf{S}(f_k)^{-1}$. Where $\mathbf{S}(f)$ is slowly varying, we can pool adjacent frequencies to gain further speed at negligible accuracy cost. Partition the block frequencies into N_c disjoint subsets $\{f_1, \dots, f_{n/(2N_b)}\} = \bigcup_{h=1}^{N_c} J_h$, each containing N_h adjacent frequencies with midpoint $\bar{f}_h$. Replace the per-frequency matrix $\mathbf{Y}(f_k)$ by the within-bin sum $\mathbf{Y}_h^* = \sum_{f \in J_h} \mathbf{Y}(f)$.

Under the assumption that $\mathbf{S}(f) \approx \mathbf{S}(\bar{f}_h)$ across each J_h , the $\mathbf{Y}(f)$ within the bin are approximately i.i.d. $\mathcal{CW}_d(T_b \mathbf{S}(\bar{f}_h), N_b)$, and by the additivity of the complex Wishart under a common scale matrix,

$$\mathbf{Y}_h^* \sim \mathcal{CW}_d(T_b \mathbf{S}(\bar{f}_h), N_b N_h). \quad (11)$$

The approximation is only as good as the constant- $\mathbf{S}$ assumption, so bin edges must avoid sharp spectral features (peaks, lines, transfer-function nulls).

The coarse-grained likelihood then takes the same form as Eq. (10) with the $n/(2N_b)$ block frequencies replaced by N_c coarse bins and a Wishart shape parameter $N_b N_h$ (the number of fine periodograms summed per coarse bin):

$$\mathcal{L}_{b,c}(\mathbf{d}|\mathbf{S}) \propto \prod_{h=1}^{N_c} |\mathbf{S}(\bar{f}_h)|^{-N_b N_h} \exp\left(-\frac{1}{T_b} \text{tr}[\mathbf{S}(\bar{f}_h)^{-1} \mathbf{Y}_h^*]\right), \quad (12)$$

with eigendecomposition

$$\mathbf{Y}_h^* = \sum_{\nu=1}^p \mathbf{u}_{\nu}^{(h)} \mathbf{u}_{\nu}^{(h)*}, \quad (13)$$

$$\mathbf{u}_{\nu}^{(h)} = \sqrt{\lambda_{\nu}^{(h)}} \mathbf{v}_{\nu}^{(h)} \quad (14)$$

B. Parametrisation using the Cholesky Decomposition

Following Rosen and Stoffer [16], Hu and Prado [17], we parametrise the *inverse* spectral density matrix through its Cholesky decomposition. Working with $\mathbf{S}^{-1}$ rather than $\mathbf{S}$ is what makes the problem tractable: the quadratic form $\mathbf{u}^* \mathbf{S}^{-1} \mathbf{u}$ splits into a sum of p decoupled quadratic terms, which in turn factorises the likelihood into p independent regressions (one per channel) that can be sampled in parallel. We write

$$\mathbf{S}(f_k)^{-1} = \mathbf{T}_k^* \mathbf{D}_k^{-1} \mathbf{T}_k, \quad (15)$$

where $\mathbf{D}_k$ is diagonal with positive entries $\delta_{1k}^2, \delta_{2k}^2, \dots, \delta_{pk}^2$, and

$$\mathbf{T}_k = \begin{pmatrix} 1 & 0 & 0 & \dots & 0 \\ -\theta_{21}^{(k)} & 1 & 0 & \dots & 0 \\ -\theta_{31}^{(k)} & -\theta_{32}^{(k)} & 1 & \ddots & \vdots \\ \vdots & \vdots & \ddots & \ddots & 0 \\ -\theta_{p1}^{(k)} & -\theta_{p2}^{(k)} & \dots & -\theta_{p,p-1}^{(k)} & 1 \end{pmatrix} \quad (16)$$

is a $p \times p$ complex unit lower triangular matrix with $\theta_{il}^{(k)}$ representing the complex-valued off-diagonal elements for $i > l$.

This parametrisation factorises the likelihood into a product of p per-channel terms,

$$\mathcal{L}(\mathbf{d}|\mathbf{S}) \propto \prod_{j=1}^p \mathcal{L}_j(\boldsymbol{\theta}_j, \boldsymbol{\delta}_j), \quad (17)$$

one for each row of $\mathbf{T}$. Each $\mathcal{L}_j$ is a univariate Whittle regression of the j th channel of the eigenvector pseudo-observations onto the preceding $(j-1)$ channels:




FIG. 1. Construction of the blocked, coarse-grained periodogram. The n -sample timeseries is partitioned into N_b non-overlapping blocks $\{b_i\}$ of length T/N_b . Each block is Fourier-transformed independently, yielding N_b single-block periodograms $\mathbf{I}^{(i)}$ each of length $n/(2N_b)$. These are summed across blocks to form the blocked periodogram $\mathbf{Y}(f_k) = \sum_{i=1}^{N_b} \mathbf{I}^{(i)}(f_k)$ (related to the Welch estimator by $\bar{\mathbf{I}} = \mathbf{Y}/N_b$), which is then summed over consecutive groups of N_h adjacent fine frequencies to produce the final coarse-grained periodogram $\mathbf{Y}_h^*$ of length $n/(2N_b N_h)$.

$$\mathcal{L}_j(\boldsymbol{\theta}_j, \boldsymbol{\delta}_j) \propto \prod_{h=1}^{N_c} \delta_{jh}^{-2N_b N_h} \exp \left(- \frac{\sum_{\nu=1}^p \left| u_{j\nu}^{(h)} - \sum_{l=1}^{j-1} \theta_{jl}^{(h)} u_{l\nu}^{(h)} \right|^2}{T_b \delta_{jh}^2} \right), \quad (18)$$

with regression coefficients $\theta_{jl}^{(h)}$ (the off-diagonal Cholesky entries at coarse bin h) and noise variance δ_{jh}^2 (the j th diagonal entry of $\mathbf{D}_h$). The eigenvectors $\mathbf{u}_\nu^{(h)}$ are those of the coarse-grained statistic $\mathbf{Y}_h^*$, evaluated

at the bin midpoint $\bar{f}_h$. The parameter blocks $\boldsymbol{\theta}_j$ and $\boldsymbol{\delta}_j$ collect all $\theta_{jl}^{(h)}$ ($l < j$) and δ_{jh} across the coarse bins $h = 1, \dots, N_c$.

In particular, for $p = 3$ channels (as in the LISA case), the three factors depend on the parameter subvectors $\boldsymbol{\delta}_1$, $(\boldsymbol{\delta}_2, \boldsymbol{\theta}_2)$, and $(\boldsymbol{\delta}_3, \boldsymbol{\theta}_3)$, respectively, as follows:

$$\mathcal{L}_1(\boldsymbol{\delta}_1) \propto \prod_{h=1}^{N_c} \delta_{1h}^{-2N_b N_h} \exp \left(- \frac{\sum_{\nu=1}^p \left| u_{1\nu}^{(h)} \right|^2}{T_b \delta_{1h}^2} \right), \quad (19)$$

$$\mathcal{L}_2(\boldsymbol{\theta}_2, \boldsymbol{\delta}_2) \propto \prod_{h=1}^{N_c} \delta_{2h}^{-2N_b N_h} \exp \left(- \frac{\sum_{\nu=1}^p \left| u_{2\nu}^{(h)} - \theta_{21}^{(h)} u_{1\nu}^{(h)} \right|^2}{T_b \delta_{2h}^2} \right), \quad (20)$$

$$\mathcal{L}_3(\boldsymbol{\theta}_3, \boldsymbol{\delta}_3) \propto \prod_{h=1}^{N_c} \delta_{3h}^{-2N_b N_h} \exp \left(- \frac{\sum_{\nu=1}^p \left| u_{3\nu}^{(h)} - \theta_{31}^{(h)} u_{1\nu}^{(h)} - \theta_{32}^{(h)} u_{2\nu}^{(h)} \right|^2}{T_b \delta_{3h}^2} \right). \quad (21)$$

² We write $\mathbf{A} \sim \mathcal{CW}_d(\mathbf{M}, m)$ if the $d \times d$ Hermitian positive definite matrix $\mathbf{A}$ has density proportional to $\det(\mathbf{M})^{-m} \exp(-\text{tr}[\mathbf{M}^{-1} \mathbf{A}])$, with mean $m\mathbf{M}$.

Note that a factorised likelihood does not in general imply a factorised posterior. Here, however, the posterior also factorises because we place independent priors across

these three parameter blocks,

$$\pi(\boldsymbol{\delta}_1, \boldsymbol{\delta}_2, \boldsymbol{\theta}_2, \boldsymbol{\delta}_3, \boldsymbol{\theta}_3) = \pi(\boldsymbol{\delta}_1) \pi(\boldsymbol{\delta}_2, \boldsymbol{\theta}_2) \pi(\boldsymbol{\delta}_3, \boldsymbol{\theta}_3), \quad (22)$$

$$\begin{aligned} p(\boldsymbol{\delta}_1, \boldsymbol{\delta}_2, \boldsymbol{\theta}_2, \boldsymbol{\delta}_3, \boldsymbol{\theta}_3 | \mathbf{d}) \propto & \mathcal{L}_1(\boldsymbol{\delta}_1) \pi(\boldsymbol{\delta}_1) \times \\ & \mathcal{L}_2(\boldsymbol{\delta}_2, \boldsymbol{\theta}_2) \pi(\boldsymbol{\delta}_2, \boldsymbol{\theta}_2) \times \\ & \mathcal{L}_3(\boldsymbol{\delta}_3, \boldsymbol{\theta}_3) \pi(\boldsymbol{\delta}_3, \boldsymbol{\theta}_3). \end{aligned} \quad (23)$$

Hence Bayes' rule yields three independent posterior factors corresponding to $\boldsymbol{\delta}_1$, $(\boldsymbol{\delta}_2, \boldsymbol{\theta}_2)$, and $(\boldsymbol{\delta}_3, \boldsymbol{\theta}_3)$. This means we can run an MCMC chain for each block separately and in parallel, which in general reduces the computation time.

C. Safe-Bayes η -tempering

The Whittle likelihood is an asymptotic approximation to the exact Gaussian likelihood of $\mathbf{Z}$, and the coarse-grained form in Eq. (12) compounds this with a within-bin constant- $\mathbf{S}$ assumption. Both approximations introduce mild misspecification. A well-known consequence of likelihood misspecification is that the posterior can over-concentrate, producing credible intervals that under-cover the truth even when the posterior mode is accurate [18].

We mitigate this with the safe-Bayes correction of Grünwald and van Ommen [18]: raise the likelihood to a power $\eta \in (0, 1]$,

$$\mathcal{L}_\eta(\mathbf{d} | \mathbf{S}) \propto \mathcal{L}_{b,c}(\mathbf{d} | \mathbf{S})^\eta, \quad (24)$$

which is equivalent to scaling the Wishart shape parameter $N_b N_h \mapsto \eta N_b N_h$ in Eq. (12). The limit $\eta \rightarrow 1$ recovers the original likelihood, while $\eta \rightarrow 0$ recovers the prior. Intermediate η down-weights the (asymptotic) likelihood relative to the prior in exactly the right amount to restore nominal coverage when the approximation is imperfect.

A useful way to read this correction is through the curvature of the log-likelihood. From Eq. (12), the observed Fisher information per coarse bin scales linearly with the number of fine periodograms summed in each bin: schematically,

$$-\partial_{\theta\theta}^2 \log \mathcal{L}_{b,c}(\mathbf{d} | \mathbf{S}(\theta)) \propto N_b N_h \mathcal{I}_0(\theta), \quad (25)$$

where $\mathcal{I}_0$ is the per-bin information matrix induced by a single complex Wishart observation. Increasing either blocking or coarse-graining therefore sharpens the log-posterior by the same factor $N_b N_h$, even though the underlying time series carries no additional information. We treat η as a tunable calibration knob rather than a fundamental quantity. For the small synthetic benchmarks in Section III A ($n = 16,384$, $N_b N_h \lesssim 10^2$) the un-tempered limit $\eta = 1$ is already well-calibrated and is used throughout. The correction becomes important for the LISA application (Section III B), where

$N_b N_h$ is two to three orders of magnitude larger and the resulting log-likelihood curvature drives both posterior over-concentration and, at sufficiently large $N_b N_h$, sampler pathologies (maximum tree-depth saturation in NUTS). An η -sensitivity study on **noise4a**, reported in Appendix C, motivates the choice $\eta = 0.5$ used for the LISA results.

D. P-spline modelling

We model the frequency-dependent Cholesky parameters with penalised B-splines,

$$\log \delta_{jk}^2 = \sum_{m=1}^{K_j} B_m(f_k) w_{j,m}^{(\delta)} \quad (26)$$

$$\Re[\theta_{jl}(f_k)] = \sum_{m=1}^{K_{jl}} B_m(f_k) w_{jl,m}^{(\Re)} \quad (27)$$

$$\Im[\theta_{jl}(f_k)] = \sum_{m=1}^{K_{jl}} B_m(f_k) w_{jl,m}^{(\Im)} \quad (28)$$

where $B_m(f_k)$ are B-spline basis functions evaluated at frequency f_k , K_j and K_{jl} are the numbers of basis functions for the j -th diagonal and (j, l) off-diagonal components, respectively, and $w_{j,m}^{(\delta)}$, $w_{jl,m}^{(\Re)}$, $w_{jl,m}^{(\Im)}$ are the corresponding spline coefficients.

Knots are placed adaptively following Maturana-Russel and Meyer [8]: for each Cholesky component the absolute-valued empirical periodogram is normalised to a probability mass function over frequency, and knots are positioned at equally-spaced quantiles of its cumulative distribution. For the off-diagonal entries the real and imaginary parts are treated as separate components, with knots placed independently on $|\Re[\theta_{jl}(f_k)]|$ and $|\Im[\theta_{jl}(f_k)]|$. Knots therefore cluster where the Cholesky component varies rapidly and spread out in smooth regions, with no manual tuning.

Each component j has its own penalty matrix $\mathbf{P}_j$ with entries

$$[\mathbf{P}_j]_{lm} = \int_0^1 B_l''(t) B_m''(t) dt. \quad (29)$$

Unlike the finite-difference approximation $\mathbf{D}_2^\top \mathbf{D}_2$ used in the original P-spline formulation of Eilers and Marx [7], this integral form is exact for non-uniform knot sequences [23].

E. Bayesian Inference

The full inference workflow proceeds in four stages. First, we place hierarchical smoothing priors on the spline coefficients of every Cholesky component. Second, we combine these priors with the η -tempered, blocked,

coarse-grained Whittle likelihood (Eq. 24) using the per-channel factorisation of Eq. (17). Third, we warm-start each MCMC chain from a draw of a stochastic variational approximation to the posterior, to bypass the expensive NUTS adaptation in low-mass-curvature regions. Fourth, we sample the warm-started posterior with NUTS and assess convergence with standard chain diagnostics. The first two stages are described below, and the remaining two stages are described in the *Implementation* and *Convergence diagnostics* paragraphs that follow.

We place hierarchical priors on the spline coefficients to enforce smoothness. For each component j :

$$\mathbf{w}_j | \phi_j \sim \mathcal{N}(\mathbf{0}, (\phi_j \tilde{\mathbf{P}}_j)^{-1}) \quad (30)$$

$$\phi_j | \delta_j \sim \text{Gamma}(\alpha_\phi, \delta_j \beta_\phi) \quad (31)$$

$$\delta_j \sim \text{Gamma}(\alpha_\delta, \beta_\delta) \quad (32)$$

where $\mathbf{w}_j$ is the vector of spline coefficients and ϕ_j, δ_j are precision hyperparameters. We use $\tilde{\mathbf{P}}_j = \mathbf{P}_j + \epsilon \mathbf{I}$ with $\epsilon = 10^{-6}$ to ensure the prior precision is strictly positive definite. Unless stated otherwise, the results below use $\alpha_\phi = \beta_\phi = 1$ and $\alpha_\delta = \beta_\delta = 1$.

Combining the η -tempered likelihood (Eq. 24) with these priors and the Cholesky factorisation of $\mathbf{S}^{-1}$, the posterior is sampled using the No-U-Turn Sampler [NUTS; 20]. The Cholesky parametrisation guarantees Hermitian positive definiteness of $\hat{\mathbf{S}}(f)$ at every frequency by construction.

a. Computational complexity. The B-spline basis matrices are evaluated once at the start of a run, after the knot locations are fixed, and reused for every gradient evaluation. Each call to $\log \mathcal{L}_\eta$ then costs $\mathcal{O}(p^2 N_c K)$ floating-point operations (using a common basis size $K = K_j = K_{jl}$ for all components): N_c coarse bins, K basis functions per component, and p^2 spline components — the p diagonal log-variances $\log \delta_{jh}^2$ plus the real and imaginary parts of the $p(p-1)/2$ off-diagonal Cholesky entries θ_{jl} . Memory is dominated by the same $p^2 N_c K$ basis-matrix storage. In practice this per-evaluation cost is small compared to the fixed overheads of NUTS warm-up and SVI initialisation, so wall-clock runtime is dominated not by the gradient-evaluation cost but by the number of leapfrog integrator steps required to traverse the posterior. As a rough scale, a single evaluation of $\log \mathcal{L}_\eta$ on the 1-year LISA configuration ($p = 3$, $N_c = 1024$, $K = 100$) takes a fraction of a millisecond on a single CPU core. Empirical end-to-end wall-clock numbers are reported alongside the tables in Section III.

b. Implementation. We implement the model in JAX [24] with NumPyro [25], exploiting automatic differentiation and JIT compilation for gradient-based sampling. Throughout we use degree-2 B-splines with second-order difference penalties. Positive precision parameters ϕ_j and δ_j are sampled on the log scale to keep the sampler away from the boundary at zero, and the frequency grid is standardised to $[0, 1]$ for numerical sta-

bility of the B-spline basis evaluation. To initialise the spline parameters we first run stochastic variational inference [SVI; 19] with a low-rank Gaussian guide (rank 16), optimised with Adam [26] at learning rate 10^{-3} for 2000 iterations. Four NUTS chains are then initialised at independent draws from the resulting variational posterior. The NUTS mass matrix is adapted during 2000 warm-up steps, followed by an additional 2000 sampling steps per chain. These sampler settings are used for both the simulation study (Section III A) and the LISA application (Section III B). Only the basis size K , the number of blocks N_b , and the tempering η vary between the two.

c. Convergence diagnostics. We assess convergence by checking that the chains satisfy $\hat{R} < 1.01$ (the Rank-normalized Gelman–Rubin potential scale reduction factor [27, 28], that the bulk and tail effective sample sizes both exceed 400, that the energy-Bayesian fraction of missing information satisfies E-BFMI > 0.3 , and that no chain shows divergent transitions or maximum tree-depth saturation. All results in Section III satisfy these criteria unless explicitly noted otherwise.

III. APPLICATION

A. Simulation study

We summarise spectral estimation accuracy with two complementary metrics. The first is the relative integrated absolute error (RIAE),

$$\text{RIAE} = \frac{\int \|\hat{S}(f) - S(f)\|_F df}{\int \|S(f)\|_F df}, \quad (33)$$

where $\|\cdot\|_F$ denotes the Frobenius norm and the integrals are evaluated numerically using Simpson’s rule. This metric measures the bias of the posterior-median estimator $\hat{S}(f)$ relative to the true spectrum $S(f)$. The second metric is the empirical coverage at level α , defined as the fraction of (frequency, matrix-element) pairs for which the true values lie inside the pointwise α credible band, averaged over realisations. This metric measures the calibration of the credible intervals. We use $\alpha = 0.9$ throughout.

We simulate from a three-dimensional VAR(2) process,

$$\mathbf{X}_t = A_1 \mathbf{X}_{t-1} + A_2 \mathbf{X}_{t-2} + \boldsymbol{\varepsilon}_t, \quad \boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \Sigma), \quad (34)$$

with

$$A_1 = \text{diag}(0.4, 0.3, 0.2), \quad (35)$$

$$A_2 = \begin{bmatrix} -0.2 & 0.5 & 0.0 \\ 0.4 & -0.1 & 0.0 \\ 0.0 & 0.0 & -0.1 \end{bmatrix}, \quad (36)$$

$$\Sigma = \begin{bmatrix} 0.25 & 0.00 & 0.08 \\ 0.00 & 0.25 & 0.08 \\ 0.08 & 0.08 & 0.25 \end{bmatrix}. \quad (37)$$

These coefficient matrices define a stationary process: the companion matrix has spectral radius $0.774 < 1$.

The one-sided theoretical spectral density matrix is

$$\mathbf{S}(f) = \frac{2}{f_s} \mathbf{H}(f) \boldsymbol{\Sigma} \mathbf{H}(f)^*, \quad \text{with} \quad (38)$$

$$\mathbf{H}(f) = \left(\mathbf{I} - \sum_{\ell=1}^2 \mathbf{A}_\ell e^{-i2\pi f \ell / f_s} \right)^{-1}. \quad (39)$$

We generate ~ 500 independent realisations of length $n = 16,384$ (with a 512-point burn-in) at sampling frequency $f_s = 1$. Each realisation is fit with $K = 10$ basis functions for each Cholesky component, while sweeping $N_b \in \{1, 2, 4, 8, 16\}$ and $N_h \in \{\text{off}, 2, 4, 8, 16, 32\}$. Safe-Bayes tempering (Section II C) is left at $\eta = 1$ throughout, with $N_h = \text{off}$ denoting no coarse-graining. A rectangular window (ENBW = 1) is applied to each block before computing the DFT. All other settings follow Section II E.

Figure 2 shows the full 3×3 spectral matrix recovered for a representative $n = 16,384$ realisation, with the NUTS and SVI posteriors overlaid. Both samplers recover the diagonal PSDs and off-diagonal cross-spectra with substantially smoother posterior medians than the raw periodogram, and the two posterior medians track each other closely (RIAE = 0.114 for NUTS versus 0.102 for SVI). The SVI credible bands, however, are visibly narrower than those of NUTS: for this realisation the empirical 90% coverage drops from 88% under NUTS to 83% under SVI. This is the well-known under-dispersion of mean-field/low-rank Gaussian variational approximations, and it is why we use SVI only as a fast NUTS initialiser and retain NUTS as the production sampler when calibrated credible bands are required. That said, the SVI fit itself completes in a few seconds and recovers an accurate posterior median, which is attractive when only a point estimate of the spectrum is needed.

Table I reports a representative slice of the (N_b, N_h) grid. Coverage is near-nominal (89.8–90.8%) across every configuration, and the mean RIAE varies by less than 2% across the grid. The number of retained frequency bins N_c falls by a factor of up to 64 between the finest ($N_b = 1$, no coarse-graining) and coarsest ($N_b = 16$, $N_h = 2$) configurations, with no measurable loss of accuracy or calibration. MCMC convergence is satisfactory across all cells: $\hat{R} \leq 1.01$ for all parameters, negligible divergent transitions, and E-BFMI > 0.6 .

Wall-clock runtime is essentially constant across the entire grid (8.7–9.6 s on a 4-core CPU node) despite the 32-fold variation in N_c . At this problem size ($K = 10$, $p = 3$, $n = 16,384$) the per-gradient cost is small enough that runtime is dominated by fixed NUTS warm-up, JIT compilation, and SVI initialisation overheads rather than by the $\mathcal{O}(p^2 N_c K)$ likelihood evaluation. The benefit of coarse-graining at this scale is therefore primarily statistical (a smaller effective $N_b N_h$ makes the posterior easier to sample) rather than computational. Coarse-graining begins to pay off in wall-clock terms once N_c or K is




FIG. 2. Representative 3×3 spectral density matrix estimate for a single $n = 16,384$ VAR(2) realisation (NUTS: RIAE = 0.114, 90% CI coverage 88%, SVI: RIAE = 0.102, 90% CI coverage 83%). Each panel shows the raw periodogram (grey), the theoretical spectrum (dotted black), and posterior medians with 90% credible bands for NUTS (blue) and SVI (orange). Diagonal panels (S_{11} , S_{22} , S_{33}) show the power spectral densities, while the off-diagonal panels show the real and imaginary parts of the cross-spectra on a linear vertical scale (the VAR(2) PSDs span less than one decade, so all panels share a linear scale).

large enough that the per-gradient cost exceeds the fixed overheads, as in the LISA application below.

B. LISA

We apply the multivariate P-spline model to the three second-generation time-delay interferometry (TDI) Michelson channels X_2 , Y_2 , Z_2 from a publicly released LISA TDI dataset [29, 30]. The data were generated with `lisainstrument` [31] under the LDC Spritz noise model [32] at a sampling rate of $f_s = 0.5$ Hz and constant 8.3 s light travel times. The six inter-spacecraft interferometer (ISI) and test-mass (TM) link measurements were propagated through `pytdi` [33] to produce the TDI 2.0 Michelson combinations used here. Two noise configurations are released:

- **noise4a** — symmetric realisation: all six MOSAs are assigned identical OMS (ISI carrier) and test-mass acceleration amplitude spectral densities equal to the LDC baseline values ($2.4 \times 10^{-15} \text{ m s}^{-2}/\sqrt{\text{Hz}}$ and $7.9 \times 10^{-12} \text{ m}/\sqrt{\text{Hz}}$, respectively). In this regime the 3×3 cross-spectral matrix is close to degenerate: $S_{XX} \approx S_{YY} \approx S_{ZZ}$, with off-diagonals equal in magnitude. A closed-form reference is available from the analytic TDI

TABLE I. 3D VAR(2) simulation results at $n = 16,384$, $K = 10$, degree-2 splines (mean $\pm$ std over ~ 500 realisations), all runs at $\eta = 1$. Coverage is the fraction of (frequency, matrix-element) pairs inside the pointwise 90% credible interval. RIAE is the Frobenius-norm relative integrated absolute error. ESS is the median effective sample size across parameters. $N_h = \text{off}$ denotes no coarse-graining ($N_h = 1$).

N_b	N_h	N_c	Coverage	RIAE	ESS	Runtime (s)
1	off	8192	0.900 ± 0.031	0.111 ± 0.004	$21,019 \pm 2726$	9.61 ± 1.76
1	2	4096	0.898 ± 0.031	0.112 ± 0.005	$24,057 \pm 2807$	9.21 ± 1.69
1	4	2048	0.900 ± 0.030	0.112 ± 0.005	$26,352 \pm 3081$	9.03 ± 1.56
1	8	1024	0.904 ± 0.029	0.111 ± 0.005	$26,640 \pm 3303$	8.92 ± 1.65
1	16	512	0.908 ± 0.029	0.110 ± 0.004	$25,480 \pm 3525$	8.81 ± 1.62
1	32	256	0.907 ± 0.030	0.109 ± 0.004	$24,011 \pm 3500$	8.82 ± 1.70
2	off	4096	0.898 ± 0.031	0.112 ± 0.005	$23,938 \pm 2939$	9.11 ± 1.63
2	2	2048	0.901 ± 0.030	0.112 ± 0.005	$26,279 \pm 2979$	9.03 ± 1.72
2	4	1024	0.904 ± 0.030	0.111 ± 0.005	$26,405 \pm 3361$	8.89 ± 1.67
2	8	512	0.907 ± 0.030	0.110 ± 0.004	$25,726 \pm 3461$	8.76 ± 1.75
2	16	256	0.907 ± 0.031	0.109 ± 0.004	$23,997 \pm 3787$	8.81 ± 1.69
4	off	2048	0.901 ± 0.030	0.112 ± 0.004	$26,053 \pm 2906$	9.05 ± 1.71
4	2	1024	0.904 ± 0.029	0.111 ± 0.005	$26,695 \pm 3221$	8.84 ± 1.64
4	4	512	0.907 ± 0.030	0.110 ± 0.004	$25,753 \pm 3636$	8.75 ± 1.72
4	8	256	0.907 ± 0.030	0.109 ± 0.004	$24,036 \pm 3622$	8.78 ± 1.68
8	off	1024	0.904 ± 0.030	0.112 ± 0.005	$26,336 \pm 3190$	8.89 ± 1.65
8	2	512	0.907 ± 0.030	0.111 ± 0.005	$25,587 \pm 3658$	8.78 ± 1.70
8	4	256	0.907 ± 0.031	0.110 ± 0.004	$24,000 \pm 3787$	8.79 ± 1.70
16	off	512	0.905 ± 0.029	0.111 ± 0.004	$25,504 \pm 3431$	8.74 ± 1.55
16	2	256	0.907 ± 0.030	0.110 ± 0.004	$23,862 \pm 3724$	8.84 ± 1.68

transfer of the link-level ASDs.

- **noise5a** — asymmetric realisation: the six per-MOSA OMS and test-mass ASDs are drawn independently from $\mathcal{U}(0.5, 2.0) \times \text{baseline}$ with a fixed seed. The resulting cross-spectral matrix has distinct diagonals and non-trivial off-diagonal structure, representing the physically relevant regime where mis-modelled cross-spectral correlations would degrade sensitivity to stochastic backgrounds. A closed-form 3×3 reference is built with the **SEGWO** package [34] by feeding the twelve measured per-MOSA ASDs through analytic OMS/TM filter shapes and projecting through PyTDI’s TDI 2.0 Michelson (X_2, Y_2, Z_2) combinations to obtain the full 3×3 XYZ covariance at every frequency.

We fit each dataset under two competing hypotheses for the 3×3 noise cross-spectrum, designed to test whether the XYZ $\rightarrow$ AET orthogonal transformation diagonalises the LISA noise covariance to within statistical resolution:

- $\mathcal{H}_{XYZ}$ — the cross-spectral matrix has fully populated off-diagonals. The multivariate Cholesky model is fit directly to the $X_2 Y_2 Z_2$ time series with $K = 100$ basis functions for *every* Cholesky component (diagonal $\log \delta_j^2$ and off-diagonal θ_{jk}). This is the unrestricted model.
- $\mathcal{H}_{AET}$ — the cross-spectral matrix is (approximately) diagonal in the A, E, T basis. The data are first rotated into AET using the orthonormal LISA TDI transformation matrix M_{AET} (whose

unit determinant introduces no log-likelihood Jacobian correction). The same multivariate Cholesky model is then fit with $K = 100$ knots on the diagonals but only $K_\theta = 10$ knots on the off-diagonals. The restricted off-diagonal prior keeps S_{AET} near diagonal. The posterior spectral matrix is rotated back to the XYZ basis for diagnostics.

For **noise4a** the AET basis diagonalises the noise exactly, so $\mathcal{H}_{\text{AET}}$ should be a valid description. For **noise5a** the per-MOSA asymmetry introduces genuine off-diagonal power in the AET basis, and $\mathcal{H}_{\text{AET}}$ is expected to be mis-specified. $\mathcal{H}_{XYZ}$ is unrestricted in both cases and serves as the reference.

For both models, knots are placed by the adaptive quantile rule described in Section II so that more knots fall near the steep low-frequency rise. We choose N_b to give blocks of order one week in length, and a fixed safe-Bayes tempering $\eta = 0.5$ (justified by the sensitivity sweep in Appendix C). Each block is tapered with a Tukey window (shape parameter $\alpha = 0.1$, ENBW ≈ 1.04) before computing the DFT to suppress spectral leakage. The three TDI transfer-function nulls at 0.030, 0.060, and 0.090 Hz are excised with half-width 1 mHz prior to fitting. All other settings follow Section II E.

Table II summarises convergence and CI-width diagnostics for the two models, two datasets, and three observation durations. All twelve LISA runs satisfy the convergence diagnostics of Section II E: $\hat{R} < 1.01$ for all parameters, bulk and tail ESS exceeding 400, E-BFMI > 0.3 , and no problematic divergent transitions or tree-depth saturation. The results split sharply along the AET-diagonalisation hypothesis.

Symmetric noise (noise4a). Both models recover the analytic spectral matrix at the half-percent level ($\text{RIAE} \approx 4.5 - 5.3 \times 10^{-3}$) across all durations. The restricted $\mathcal{H}_{AET}$ model is no worse than $\mathcal{H}_{XYZ}$ in posterior-median accuracy, consistent with the AET basis exactly diagonalising the equal-link noise.

Asymmetric noise (noise5a). The two models diverge by more than an order of magnitude. $\mathcal{H}_{AET}$ stalls at $\text{RIAE} \approx 1.8 - 2.2 \times 10^{-2}$ at all durations — a floor set by the residual off-diagonal AET power that the restricted prior cannot represent. $\mathcal{H}_{XYZ}$, in contrast, reaches $\text{RIAE} \approx 9 - 11 \times 10^{-4}$, about $20\times$ smaller, and continues to improve with duration. This is the central diagnostic of the LISA study: when the per-MOSA noise levels are asymmetric (the physically realistic regime) the AET rotation does *not* diagonalise the cross-spectral matrix, and a multivariate model that estimates the full Cholesky factor is required.

Posterior uncertainty contracts as expected for both models. The median relative CI width on the diagonal PSDs falls from $\approx 4.7\%$ at 1 month to $\approx 1.2\%$ at 1 year, a reduction of roughly $4\times$, consistent with the $\sqrt{T}$ scaling expected for variance-dominated Whittle inference, and is essentially identical across models and datasets, confirming that the per-channel resolution is set by the data and not by the choice of basis.

Figure 3 overlays the three durations for both datasets under the $\mathcal{H}_{XYZ}$ model. The diagonal PSDs are recovered with sub-percent bias and CI widths that contract with $\sqrt{T}$, while the off-diagonal coherences (the cross-channel quantities most directly tied to stochastic-background and parameter-estimation pipelines [35]) are resolved at the few-percent level by 6 months and tightened by a further factor of ~ 1.7 at 1 year. A localised feature visible in all three off-diagonal panels of Fig. 3 is a cluster of small-amplitude wiggles in the posterior-median coherence near the TDI transfer-function null at ≈ 0.06 Hz (and, more weakly, near 0.03 and 0.09 Hz). These arise because the diagonal PSDs $S_{ii}(f)$ drop by several orders of magnitude across each null, and the P-spline basis — which is smooth on a log scale but agnostic to the location and shape of the dips — has to absorb this dynamic range from the data alone. Small residual misfits at the bottoms of the dips are amplified in the coherence, producing the visible ripple. Incorporating the analytic TDI transfer functions directly into the prior mean (rather than only excising a 1 mHz window around each null) may be a natural way to suppress this artefact and is left for future work. Wall-clock runtime grows much more slowly than the data volume: a $12\times$ increase in mission duration (1 mo $\rightarrow$ 1 yr) costs only a $\sim 7\times$ increase in wall time for $\mathcal{H}_{XYZ}$. The 1-month and 6-month fits (~ 120 – 245 s) are dominated by fixed NUTS warm-up and SVI overhead, while the 1-year fit (~ 850 – 1090 s) begins to expose the per-gradient $\mathcal{O}(p N_b N_c K)$ cost as N_b reaches 52. Even the 1-year fit completes in under twenty minutes on a single 4-core CPU node, roughly a $2\times$ speed-up at 1 month and $1.5\times$ at 6 months over the

previous-generation single-model pipeline, with comparable throughput at 1 year. The $\mathcal{H}_{AET}$ model is consistently faster than $\mathcal{H}_{XYZ}$ at fixed duration (e.g. 632 s vs. 853 s at 1 yr on **noise4a**), reflecting the smaller off-diagonal basis, but at the cost of the order-of-magnitude RIAE inflation on the asymmetric **noise5a** regime.

IV. DISCUSSION

We have presented a multivariate P-spline framework for estimating the cross-spectral density matrix of stationary vector time series. The Cholesky parametrisation of $\mathbf{S}^{-1}$ guarantees Hermitian positive definiteness at every frequency and factorises the likelihood into p independent per-channel regressions that can be sampled in parallel. A blocked, coarse-grained Whittle likelihood combined with safe-Bayes η -tempering yields near-nominal 90% credible-interval coverage (≈ 0.91) across a wide (N_b, N_h) grid on the 3D VAR(2) benchmark, while reducing the number of inference-time frequency bins by up to two orders of magnitude.

On the synthetic 3D VAR(2) benchmark the untempered estimator ($\eta = 1$) achieves coverage ≈ 0.91 at every (N_b, N_h) configuration tested. On the LISA TDI data the same untempered estimator over-concentrates. A fixed safe-Bayes tempering $\eta = 0.5$ recovers credible bands that contract with $\sqrt{T}$ and remain well-behaved across all three observation durations (Appendix C).

a. η -tempering and debiased Whittle. The η -sweep of Appendix C shows that coverage is a strong, monotone function of η at fixed (N_b, N_h, K) , but does not yet identify a principled rule for choosing η across regimes. Mapping the optimal η as a function of $N_b N_h$, K , and the underlying spectral curvature is the most pressing methodological question raised by this work. A complementary direction is to replace η -tempering with a debiased Whittle likelihood [36], which corrects the finite-sample bias of the periodogram analytically rather than down-weighting the entire likelihood. A debiased formulation may reduce the reliance on η -tempering by correcting finite-sample bias analytically, though steep likelihood curvature from high $N_b N_h$ would likely persist and still require attenuation.

b. Informative priors. The hierarchical Gaussian smoothing prior used here is near-uninformative: each weight has marginal variance set by a high-level hyperparameter (ϕ_j, δ_j) with no scale imposed by the data ahead of inference. In practice this places no bound on where the spline curves can sit, and most of NUTS warm-up is spent locating a region of weight space consistent with even a coarse Welch estimate. A more informative prior, centred on a fast initial estimate of $\log \delta_j^2(f)$ and the off-diagonal Cholesky terms (e.g. from a Welch-and-smooth pre-pass), with ϕ_j controlling allowed deviations from that estimate, would (i) make the smoothing hyperparameters physically meaningful as fractional excursion bounds, (ii) shorten warm-up substantially, and (iii) re-

TABLE II. LISA TDI $X_2Y_2Z_2$ results at fixed safe-Bayes tempering $\eta = 0.5$ for the symmetric (**noise4a**) and asymmetric (**noise5a**) simulations, fit with the unrestricted multivariate XYZ model ($\mathcal{H}_{XYZ}$) and the AET-rotated model with restricted off-diagonal basis ($\mathcal{H}_{AET}$, $K_\theta = 10$), at three observation durations. N_b is the number of Wishart time blocks (chosen to give ~ 7 -day blocks: $\lfloor T/(7\text{d}) \rfloor$) and N_c is set to 1024. RIAE is the Frobenius-norm relative integrated absolute error of the posterior-median spectral matrix against the analytic reference (closed-form equal-link LDC for **noise4a**, SEGWO with measured per-MOSA ASDs for **noise5a**). Δ_{PSD} is the median over frequencies of the pointwise relative width $[\hat{S}^{95\%}(f) - \hat{S}^{5\%}(f)]/\hat{S}^{50\%}(f)$ of the 90% credible band on the diagonal PSDs (computed in the XYZ basis for both models). Runtime is wall-clock seconds for one full SVI+NUTS pipeline run on a single 4-core CPU node.

Noise	Duration	N_b	$\mathcal{H}_{AET}$			$\mathcal{H}_{XYZ}$		
			RIAE	Δ_{PSD}	t (s)	RIAE	Δ_{PSD}	t (s)
4a	1 mo	4	4.5×10^{-3}	4.8%	117	4.7×10^{-3}	4.6%	119
	6 mo	25	4.7×10^{-3}	1.9%	224	5.3×10^{-3}	1.9%	181
	1 yr	52	4.4×10^{-3}	1.2%	632	5.0×10^{-3}	1.2%	853
5a	1 mo	4	2.2×10^{-2}	4.9%	119	1.1×10^{-3}	4.5%	177
	6 mo	25	1.8×10^{-2}	1.9%	218	9.8×10^{-4}	1.9%	245
	1 yr	52	2.2×10^{-2}	1.3%	728	8.6×10^{-4}	1.3%	1089




FIG. 3. Posterior summaries of the LISA TDI $X_2Y_2Z_2$ spectral matrix from the unrestricted $\mathcal{H}_{XYZ}$ model for the symmetric **noise4a** (left) and asymmetric **noise5a** (right) datasets at safe-Bayes tempering $\eta = 0.5$, shown at three observation durations: 1 month (blue), 6 months (orange), and 12 months (green). Diagonal panels show the posterior 90% credible bands for the diagonal PSDs S_{XX} , S_{YY} , S_{ZZ} . Lower-triangle panels show the coherence $|C_{ij}(f)|$ (Eq. 4) for each pair. Welch reference (grey) is overplotted.

duce the prior–data contention that drives over-coverage at small η in the appendix sweep. We expect this to give the largest practical improvement for production LISA pipelines.

c. Non-stationarity. The Whittle likelihood assumes wide-sense stationarity. Both **noise4a** and **noise5a** are stationary by construction: realisations of an LDC noise model with constant per-MOSA amplitude spectral densities and constant 8.3s light travel times. The real LISA noise environment is non-stationary on multiple timescales — constellation breathing modulates the arm lengths over the year-long orbital period, and per-link noise levels drift on timescales from hours (thermal cycling, glitches) to months (component aging). A natural and tractable extension is to assume *piecewise*

stationarity: partition the mission timeline into segments short enough that within-segment stationarity is a good approximation but long enough that the Whittle approximation remains accurate, and either fit each segment independently or couple consecutive segments through a smoothness prior on the frequency-dependent Cholesky parameters in the time direction. This would extend the present model to a time-frequency spectral matrix $\mathbf{S}(t, f)$ at the cost of one additional smoothing dimension.

d. Improved variational families. As reported in Section III A, the current low-rank ($r = 16$) multivariate Gaussian SVI guide recovers an accurate posterior median but produces visibly narrower credible bands than NUTS, with a corresponding drop in empirical coverage. It is also too restrictive to represent the heavy tails or

weak multi-modality occasionally seen in the off-diagonal Cholesky weights at small η . NUTS is therefore retained as the production sampler in this work, and SVI is used only as an initialiser. Because the model is implemented in NumPyro, the variational family is a swappable component: replacing the Gaussian guide with a normalising flow (e.g. a neural spline flow or a masked autoregressive / inverse autoregressive flow) is essentially a one-line change. A systematic study of which variational families recover calibrated PSD posteriors — so that SVI could serve as a stand-alone approximation at a fraction of the NUTS wall-clock cost, attractive for the routine, parameter-estimation-grade PSD updates expected over the LISA mission lifetime — is left for future work, as is a corresponding ablation of NUTS with and without the SVI warm-start (quantifying its effect on warm-up length, divergences, and time-to-convergence). A further speculative direction is to use the variational posterior not only to initialise NUTS but to recentre the spline-weight priors on the variational mean, with the variational scale informing the prior width. The trade-off is that this makes the prior data-dependent and informative, which is undesirable when calibration of the final posterior is the goal. We flag it as a direction worth exploring rather than recommend it.

e. Other extensions. The within-bin constant- $\mathbf{S}$ approximation underlying coarse-graining (Eq. 12) can bias estimates near sharp spectral features. Adaptive bin widths that narrow around lines and transfer-function nulls are a natural extension. Hybrid parametric/nonparametric models that embed known instrumental features within the P-spline envelope are another route to improving sharp-feature accuracy without sacrificing the smoothness elsewhere.

Data and Software Availability

The software developed for this project is open-source and publicly available from the GitHub repository <https://github.com/nz-gravity/LogPsplinePSD>, which contains all source code, example scripts, and configuration files needed to reproduce the results. Installation instructions and dependency information (tested on Python ≥ 3.10) are provided in the repository README. The software is released under the MIT License, permitting free use, modification, and distribution. LISA time series data is available at [29, 30]. Scripts for the LISA analysis are provided from the GitHub repository http://github.com/nz-gravity/lisa_multivar_logpsplines_study.git

ACKNOWLEDGMENTS

We thank Martina Muratore, Federico Pozzoli, Nikolaos Karnesis, Quentin Baghi, Jean-Baptiste Bayle, Ollie Burke, and others in the LISA Noise Non-Stationarities

Group (part of the “Deep analysis group” of the Distributed Data Processing Centre, DDPC) for helpful discussions. We additionally thank Jean-Baptiste Bayle for providing the LISA dataset, and the University of Glasgow for the computing resources that supported JB in simulating the dataset. We acknowledge the use of the SEGWO package [34], developed by J.-B. Bayle and O. Hartwig, used here to construct the analytic 3×3 cross-spectral reference for the asymmetric `noise5a` realisation. PMR, RM, and AV gratefully acknowledge support from the Marsden Fund Council grants MFP-UOA2131 and MFP-UOA2531, funded by the New Zealand Government and managed by the Royal Society Te Apārangi. This work was performed on the OzSTAR national facility at Swinburne University of Technology. The OzSTAR program receives funding in part from the Astronomy National Collaborative Research Infrastructure Strategy (NCRIS) allocation provided by the Australian Government, and from the Victorian Higher Education State Investment Fund (VHESIF) provided by the Victorian Government.

Appendix A: Notation

Table III lists the symbols used throughout the paper, grouped by the stage of the model in which they appear.

Appendix B: Additional bivariate VAR(2) sensitivity analyses

This appendix benchmarks the proposed estimator against the bivariate Bayesian baselines available in the literature on a controlled VAR(2) problem with closed-form ground truth. Following Liu *et al.* [21], we simulate 500 independent realisations of a bivariate time series from a vector autoregressive model of order 2, VAR(2), using three different sample sizes $n = 256, 512, 1024$ (refer to Liu *et al.* [Section 4.2, 21] for the definitions of the VAR(2) models). The spectral densities are estimated using the multivariate P-spline method with $K = 16$ basis functions and cubic splines ($d = 3$) for each spectral component. We validate the multivariate P-spline method using synthetic data from Vector Autoregressive (VAR) processes with known theoretical spectral properties. The theoretical spectral density matrix provides exact ground truth for comparison. To benchmark against existing Bayesian multivariate spectral estimators, we compare to the variational Bayes (VB) and variational nonparametric correction (VNPC) methods of Liu *et al.* [21] using the same L_2 error criterion. The L_2 error is the un-normalised counterpart of the matrix RIAE used in the main text: it is the Frobenius-norm integrated absolute error of the posterior-median spectral matrix without dividing by $\int \|S(f)\|_F df$.

Table IV summarises the L_2 errors across the 500 independent realisations. The multivariate P-spline estima-

Symbol	Meaning
<i>Time series & sampling</i>	
T, Δ_t, f_s	Total observation duration, sampling interval, sampling frequency; $T = n\Delta_t$, $\Delta_t = 1/f_s$.
f_{Ny}	Nyquist frequency, $f_{Ny} = f_s/2$.
n, p	Number of time samples per channel, number of channels.
Δ_f, f_k	Frequency resolution $\Delta_f = 1/T$; k th DFT frequency $f_k = k\Delta_f$.
N	Number of positive Fourier frequencies, $N = \lfloor n/2 \rfloor$.
<i>Blocking & coarse-graining</i>	
N_b, T_b	Number of non-overlapping time blocks; block duration $T_b = T/N_b$ (n/N_b samples per block).
$N_h, J_h, \bar{f}_h$	Number of fine frequencies aggregated per coarse bin; J_h the h th bin (consecutive subset of fine frequencies), with midpoint $\bar{f}_h$.
N_c	Number of coarse-grain bins, $N_c = n/(2N_bN_h)$.
<i>Spectral density</i>	
$\mathbf{Z}$	Observed p -channel time series, $\mathbf{Z} \in \mathbb{R}^{n \times p}$.
$\mathbf{d}(f_k)$	DFT vector, $\mathbf{d}(f_k) \in \mathbb{C}^p$.
$\Gamma(h), \gamma_{lm}(h)$	Autocovariance matrix at lag h and its (l, m) entry.
$\mathbf{S}(f_k)$	$p \times p$ Hermitian positive-definite spectral density matrix; Fourier transform of Γ .
<i>Periodogram & Wishart statistics</i>	
$\mathbf{I}^{(i)}(f_k), \bar{\mathbf{I}}(f_k)$	Block- i periodogram, $\mathbf{I}^{(i)} = \mathbf{d}^{(i)} \mathbf{d}^{(i)*} \sim \mathcal{CW}_p(T_b \mathbf{S}, 1)$; block-averaged periodogram $\bar{\mathbf{I}} = N_b^{-1} \sum_i \mathbf{I}^{(i)}$.
$\mathbf{Y}(f_k)$	Summed periodogram / Wishart statistic, $\mathbf{Y} = N_b \bar{\mathbf{I}} \sim \mathcal{CW}_p(T_b \mathbf{S}, N_b)$.
$\mathbf{Y}_h^*$	Coarse-grained statistic, $\mathbf{Y}_h^* = \sum_{f_k \in J_h} \mathbf{Y}(f_k) \sim \mathcal{CW}_p(T_b \mathbf{S}(\bar{f}_h), N_b N_h)$.
$\lambda_\nu^{(k)}, \mathbf{v}_\nu^{(k)}, \mathbf{u}_\nu^{(k)}$	ν th eigenvalue/unit eigenvector of $\mathbf{Y}(f_k)$; scaled eigenvector $\mathbf{u}_\nu^{(k)} = \sqrt{\lambda_\nu^{(k)}} \mathbf{v}_\nu^{(k)}$.
<i>Cholesky parametrisation</i>	
$\mathbf{T}_k, \mathbf{D}_k$	Factors in $\mathbf{S}(f_k)^{-1} = \mathbf{T}_k^* \mathbf{D}_k^{-1} \mathbf{T}_k$: $\mathbf{T}_k$ unit lower-triangular, $\mathbf{D}_k$ diagonal with entries $\delta_{1k}^2, \dots, \delta_{pk}^2$.
$\theta_{jl}^{(k)}, \delta_{jk}^2$	Off-diagonal entry (j, l) of $\mathbf{T}_k$ ($j > l$) and j th diagonal entry of $\mathbf{D}_k$.
$\boldsymbol{\theta}_j, \boldsymbol{\delta}_j$	Parameter vectors for channel j : collect $\theta_{jl}^{(k)}$ ($l < j$) and δ_{jk}^2 across $k = 1, \dots, n/(2N_b)$.
<i>P-spline model</i>	
$K, B_m(f_k)$	Number of B-spline basis functions (model hyperparameter, distinct from frequency index k); m th basis evaluated at f_k .
M_j, M_{jl}	Basis sizes for the j th diagonal and (j, l) off-diagonal components.
$\mathbf{w}_j^{(\delta)}, \mathbf{w}_{jl}^{(\Re)}, \mathbf{w}_{jl}^{(\Im)}$	Spline coefficient vectors for $\log \delta_{jk}^2, \Re[\theta_{jl}], \Im[\theta_{jl}]$.
$\mathbf{P}_j, \phi_j$	Penalty matrix $[\mathbf{P}_j]_{lm} = \int_0^1 B_l'' B_m'' dt$ and smoothing precision hyperparameter.

TABLE III. Notation used throughout the paper. Symbols are grouped by the stage of the model they belong to.

tor is competitive with both comparators at every sample size and attains the lowest mean error in all three settings. The mean L_2 error decreases from 0.095 at $n = 256$ to 0.072 at $n = 512$ and 0.055 at $n = 1024$, compared with 0.121, 0.091, and 0.066 for VB and 0.129, 0.103, and 0.082 for VNPC. This is the main message we want from the bivariate benchmark: the proposed Log-P-spline method achieves accuracy of the same order as established alternatives, while in these simulations being modestly more accurate on average.

TABLE IV. Bivariate VAR(2) benchmark using the L_2 error metric from Liu *et al.* [21]. Entries are mean $\pm$ standard deviation over 500 independent realisations. Smaller values are better.

n	P-spline	VB	VNPC
256	0.095 $\pm$ 0.022	0.121 $\pm$ 0.029	0.129 $\pm$ 0.033
512	0.072 $\pm$ 0.018	0.091 $\pm$ 0.027	0.103 $\pm$ 0.025
1024	0.055 $\pm$ 0.013	0.066 $\pm$ 0.016	0.082 $\pm$ 0.017




FIG. 4. Safe-Bayes η -sensitivity on the symmetric LISA `noise4a` dataset at three observation durations (1 month, 6 months, 1 year). Top: median pointwise relative width of the 90% credible band on the diagonal PSDs. Bottom: matrix RIAE (narrow y -range—RIAE varies by $\lesssim 5\%$ across the full sweep). The vertical dotted line marks $\eta = 0.5$, the value used in Section III B.

Appendix C: η -tempering sensitivity on `noise4a`

To support the choice of tempering parameter for the LISA application, we re-ran the `noise4a` dataset at $\eta \in \{0.01, 0.03, 0.1, 0.25, 0.5, 0.75, 1\}$ for each of the three observation durations (1, 6, and 12 months) used in Section III B, holding all other settings fixed. Figure 4 summarises the result.

Two features emerge from Figure 4. First, the matrix RIAE is essentially flat across the entire sweep: it varies by less than $\pm 5\%$ about its mean at every duration, confirming that the posterior-median point estimate is nearly insensitive to η . Second, credible-band width contracts monotonically with η , by roughly an order of magnitude from $\eta = 0.01$ to $\eta = 1$, with the steepest contraction at the longest duration—consistent with the per-bin Fisher information scaling as $N_b N_h$ (Eq. (25)).

The vertical dotted line marks the choice $\eta = 0.5$ used in Section III B. This value gives a practical compromise: credible bands are meaningfully wide at all three durations while the point estimate remains accurate throughout. Reaching nominal coverage would require yet smaller η or a more flexible model (e.g. adaptive coarse-graining near the TDI nulls). We leave this to future work.

-
- [1] T. B. Littenberg and N. J. Cornish, Bayesian inference for spectral estimation of gravitational wave detector noise, *Physical Review D* **91**, 084034 (2015).
 - [2] P. A.-S. et al., Laser interferometer space antenna (2017), [arXiv:1702.00786](#).
 - [3] C. Kirch, M. C. Edwards, A. Meier, and R. Meyer, Beyond Whittle: Nonparametric Correction of a Parametric Likelihood with a Focus on Bayesian Time Series Analysis, *Bayesian Analysis* **14**, 1037 (2019).
 - [4] Q. Baghi, N. Karnesis, J.-B. Bayle, M. Besançon, and H. Inchauspé, Uncovering gravitational-wave backgrounds from noises of unknown shape with lisa, *Journal of cosmology and astroparticle physics* **2023**, 66 (2023).
 - [5] J. Liu, A. Vajpeyi, R. Meyer, K. Janssens, J. E. Lee, P. Maturana-Russel, N. Christensen, and Y. Liu, Variational inference for correlated gravitational wave detector network noise, *Phys. Rev. D* **111**, 062003 (2025).
 - [6] M. C. Edwards, R. Meyer, and N. Christensen, Bayesian nonparametric spectral density estimation using b-spline priors, *Statistics and computing* **29**, 67 (2019).
 - [7] P. H. C. Eilers and B. D. Marx, Flexible smoothing with B-splines and penalties, *Statistical Science* **11**, 89 (1996).
 - [8] P. Maturana-Russel and R. Meyer, Bayesian spectral density estimation using P-splines with quantile-based knot placement, *Computational Statistics* **36**, 2055 (2021).
 - [9] N. Aimen, P. Maturana-Russel, A. Vajpeyi, N. Christensen, and R. Meyer, Bayesian power spectral density estimation for LISA noise based on penalized splines with a parametric boost, *Physical Review D* **113**, 024022 (2026).
 - [10] Q. Baghi, N. Karnesis, J.-B. Bayle, M. Besançon, and H. Inchauspé, Uncovering gravitational-wave backgrounds from noises of unknown shape with LISA, *Journal of Cosmology and Astroparticle Physics* **2023** (04), 066.
 - [11] L. Santini, M. Muratore, J. Gair, and O. Hartwig, Bayesian inference on stochastic gravitational wave backgrounds with LISA, *Physical Review D* **112**, 084050 (2025).
 - [12] M. Muratore, *Instrumental modelling and noise reduction algorithms for the Laser Interferometer Space Antenna*, *Ph.D. thesis*, Gottfried Wilhelm Leibniz Universität Han-

- nover (2021).
- [13] F. Cireddu, M. Wils, I. C. F. Wong, P. T. H. Pang, T. G. F. Li, and W. Del Pozzo, Likelihood for a network of gravitational-wave detectors with correlated noise, *Phys. Rev. D* **110**, 104060 (2024).
 - [14] O. Hartwig, M. Lilley, M. Muratore, and M. Pieroni, Stochastic gravitational wave background reconstruction for a nonequilateral and unequal-noise lisa constellation, *Phys. Rev. D* **107**, 123531 (2023).
 - [15] M. Muratore, J. Gair, and L. Speri, Impact of the noise knowledge uncertainty for the science exploitation of cosmological and astrophysical stochastic gravitational wave background with lisa, *Phys. Rev. D* **109**, 042001 (2024).
 - [16] O. Rosen and D. S. Stoffer, Automatic estimation of multivariate spectra via smoothing splines, *Biometrika* **94**, 335 (2007).
 - [17] Z. Hu and R. Prado, Fast Bayesian inference on spectral analysis of multivariate stationary time series, *Computational Statistics & Data Analysis* **178**, 107596 (2023).
 - [18] P. Grünwald and T. van Ommen, Inconsistency of Bayesian inference for misspecified linear models, and a proposal for repairing it, *Bayesian Analysis* **12**, 1069 (2017).
 - [19] M. D. Hoffman, D. M. Blei, C. Wang, and J. Paisley, Stochastic variational inference, *Journal of machine learning research* (2013).
 - [20] M. D. Hoffman, A. Gelman, *et al.*, The no-u-turn sampler: adaptively setting path lengths in hamiltonian monte carlo., *J. Mach. Learn. Res.* **15**, 1593 (2014).
 - [21] Y. Liu, C. Kirch, J. E. Lee, and R. Meyer, *A nonparametrically corrected likelihood for bayesian spectral analysis of multivariate time series* (2024).
 - [22] J. P. Grainger, A. M. Sykulski, K. Ewans, H. F. Hansen, and P. Jonathan, A multivariate pseudo-likelihood approach to estimating directional ocean wave models, *Journal of the Royal Statistical Society Series C: Applied Statistics* **72**, 544 (2023).
 - [23] M. P. Wand and J. T. Ormerod, On semiparametric regression with o’sullivan penalized splines, *Australian & New Zealand Journal of Statistics* **50**, 179 (2008).
 - [24] J. Bradbury, R. Frostig, P. Hawkins, M. J. Johnson, C. Leary, *et al.*, *JAX: composable transformations of Python+NumPy programs* (2018).
 - [25] D. Phan, N. Pradhan, and M. Jankowiak, Composable effects for flexible and accelerated probabilistic programming in pyro, *arXiv preprint arXiv:1912.11554* (2019).
 - [26] D. P. Kingma and J. Ba, Adam: A method for stochastic optimization, in *3rd International Conference on Learning Representations (ICLR)* (2015) [arXiv:1412.6980 \[cs.LG\]](#).
 - [27] A. Gelman and D. B. Rubin, Inference from Iterative Simulation Using Multiple Sequences, *Statistical Science* **7**, 457 (1992).
 - [28] R. Kumar, C. Carroll, A. Hartikainen, and O. Martin, Arviz a unified library for exploratory analysis of bayesian models in python, *Journal of Open Source Software* **4**, 1143 (2019).
 - [29] J.-B. Bayle, Lisa sgwb dataset (noise-4a), [10.5281/zenodo.15698080](#) (2025).
 - [30] J.-B. Bayle, Lisa sgwb dataset (noise-5a), [TODO](#) (2026).
 - [31] J.-B. Bayle, O. Hartwig, and M. Staab, *LISA instrument* (2024).
 - [32] M. Le Jeune, S. Babak, Q. Baghi, J.-B. Bayle, E. Castelli, and N. Korsakova, Lisa data challenge spritz (ldc2b), [10.5281/zenodo.7436568](#) (2022).
 - [33] M. Staab, J.-B. Bayle, and O. Hartwig, *PyTDI* (2022).
 - [34] J.-B. Bayle and O. Hartwig, *SEGWO: Sensitivity estimator for gravitational-wave observatories*.
 - [35] M. L. Katz *et al.*, The LISA global fit, *arXiv preprint arXiv:2404.12571* (2024).
 - [36] A. M. Sykulski, S. C. Olhede, A. P. Guillaumin, J. M. Lilly, and J. J. Early, The debiased Whittle likelihood, *Biometrika* **106**, 251 (2019).